

CMSC424: Database Design

Module: Introduction/Overview

Instructor: Amol Deshpande
amol@umd.edu

Welcome to CMSC424: Database Design

▶ Course description

- *Students are introduced to database systems and motivates the database approach as a mechanism for modeling the real world. An in-depth coverage of the relational model, logical database design, query languages, and other database concepts including query optimization, concurrency control; transaction management, and log based crash recovery. Distributed and Web database architectures are also discussed.*

▶ Basics

- Typical in-person instruction, 2 classes a week etc.
 - Will record lectures but quality may not be as good
 - Plan to substitute a portion of the in-class teaching with videos + in-class activities
- Will make available pre-recorded videos from Fall 2020
 - Expect about 75% overlap
- Use the discussion forums, office hours as much as possible – assignments are to be done individually, but you are allowed to discuss among yourselves
- Fill out the survey covering your background, your expectations, and potential issues (on Gradescope)

No Laptops without permission !

Logistics

- ▶ Instructor: Amol Deshpande
 - 5154 IRB, amol@umd.edu
 - Class Webpage: ELMS
- ▶ Email to me: write CMSC424 in the title
 - Only if Piazza not suitable for some reason
 - Piazza allows private posts
- ▶ TAs: Xincheng Yang, Fnu Ferry, Vanshika Mehta, Ayush Sood

Logistics

- ▶ Textbook:
 - Database System Concepts
 - Seventh Edition (6th Edition will be fine too)
 - Abraham Silberschatz, Henry F. Korth, S. Sudarshan
- ▶ GitHub:
 - To distribute programming assignments and Slides
 - <https://github.com/umddb/cmsc424-spring2024>
- ▶ Piazza: First resort for any questions
- ▶ ELMS: General announcements will be posted there
- ▶ GradeScope
 - We will use this for assignments, grading exams, etc.

Grading

- ▶ 3 Midterms – 30%
 - 20-25 mins at the beginning of the class
- ▶ Final (30%) – minimum of 50% on this + midterms combined to pass
- ▶ 6 Assignments (24%) + Setup (1%)
 - SQL, Advanced SQL, SQL + Programming Languages, Normalization/E-R Models (not programming), MongoDB, Spark
 - Late Policy
 - Automatic one-week extension without penalty, but with an explanation of why
 - Beyond that: submissions allowed through the end of the semester with penalties
 - Assignment-0 (Setup) due next Friday
- ▶ In-class Activities (15%)
 - Each week – will devote some time in the class to do some worksheets
 - Videos to watch before Tuesday class each week (starting next) – aim to release by Friday night

ChatGPT etc.

- ▶ Feel free to use for learning purposes
 - e.g., to debug issues with set up, or to understand some concepts better
 - Beware of the hallucination and other issues
- ▶ Will try to provide some starting points and prompts
- ▶ Not allowed for directly solving assignments, although debugging would be acceptable

CMSC424: Database Design

Module: Introduction/Overview

**Background; What
We Will Cover**

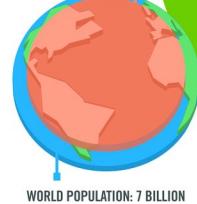
Instructor: Amol Deshpande
amol@umd.edu

Motivation

- ▶ Book Chapters (6th Edition)
 - 1.1, 1.2, 1.3, 1.4, 1.9 (to some extent)
- ▶ Key Topics
 - Why managing large volumes of data is difficult
 - Database Systems and Architectures
 - Data Models and Data Abstraction
 - What we will cover in this course

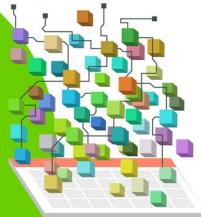
40 ZETTABYTES
[43 TRILLION GIGABYTES]
of data will be created by
2020, an increase of 300
times from 2005

**6 BILLION
PEOPLE**
have cell phones



Volume SCALE OF DATA

It's estimated that
2.5 QUINTILLION BYTES
[2.3 TRILLION GIGABYTES]
of data are created each day



Most companies in the
U.S. have at least
100 TERABYTES
[100,000 GIGABYTES]
of data stored

The New York Stock Exchange
captures
**1 TB OF TRADE
INFORMATION**
during each trading session



Velocity ANALYSIS OF STREAMING DATA

Modern cars have close to
100 SENSORS
that monitor items such as
fuel level and tire pressure

By 2016, it is projected
there will be
**18.9 BILLION
NETWORK
CONNECTIONS**
- almost 2.5 connections
per person on earth



The FOUR V's of Big Data

From traffic patterns and music downloads to web history and medical records, data is recorded, stored, and analyzed to enable the technology and services that the world relies on every day. But what exactly is big data, and how can these massive amounts of data be used?

As a leader in the sector, IBM data scientists break big data into four dimensions: **Volume**, **Velocity**, **Variety** and **Veracity**

Depending on the industry and organization, big data encompasses information from multiple internal and external sources such as transactions, social media, enterprise content, sensors and mobile devices. Companies can leverage data to adapt their products and services to better meet customer needs, optimize operations and infrastructure, and find new sources of revenue.

By 2015
4.4 MILLION IT JOBS
will be created globally to support big data,
with 1.9 million in the United States



As of 2011, the global size of data in healthcare was
estimated to be

150 EXABYTES
[161 BILLION GIGABYTES]



**30 BILLION
PIECES OF CONTENT**

are shared on Facebook
every month



Variety DIFFERENT FORMS OF DATA

By 2014, it's anticipated
there will be

**420 MILLION
WEARABLE, WIRELESS
HEALTH MONITORS**

**4 BILLION+
HOURS OF VIDEO**
are watched on
YouTube each month



400 MILLION TWEETS
are sent per day by about 200
million monthly active users

Poor data quality costs the US
economy around

\$3.1 TRILLION A YEAR



**1 IN 3 BUSINESS
LEADERS**

don't trust the information
they use to make decisions



**27% OF
RESPONDENTS**

in one survey were unsure of
how much of their data was
inaccurate

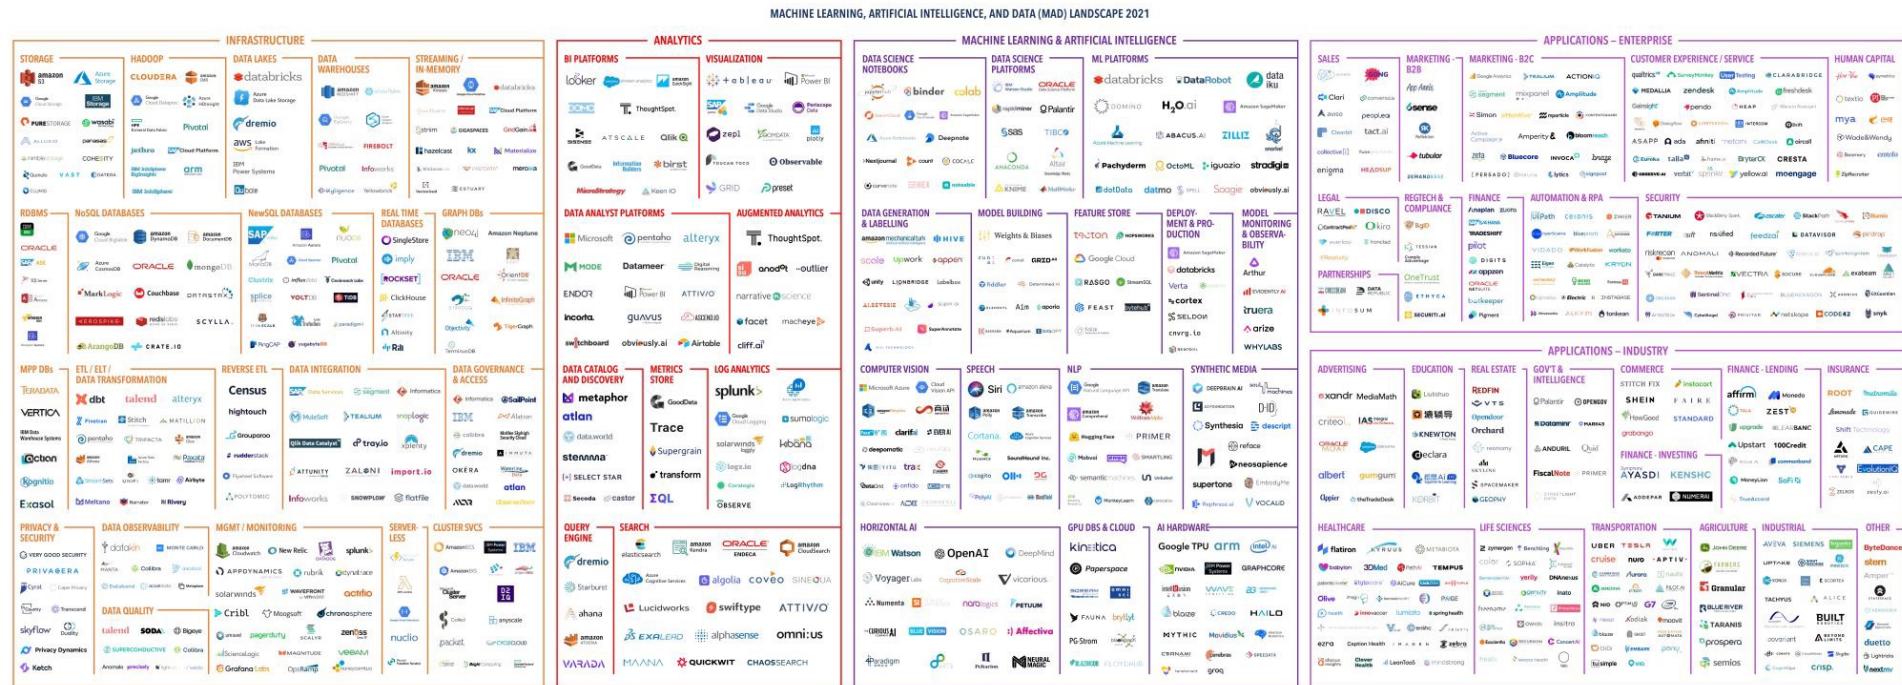
Veracity UNCERTAINTY OF DATA

IBM

Data Management Systems

Data management systems at the center of most of the new innovative technologies

An incredible amount of ongoing work in building new types of systems



INFRASTRUCTURE



Data Management

A large fraction of the data still in traditional DBMS systems

Still open and active research areas about improving performance, energy efficiency, new functionalities, changing hardware spectrum (SSDs) and so on...

Much of the data not stored in traditional database systems today

For a variety of fairly valid reasons

- Stream processing systems (focusing on *streaming* data)
- Special-purpose data warehousing systems (most start from some RDBMS)
- Batch analysis frameworks (like Hadoop, Pregel, Spark, ...)

Typically data stored in distributed file systems

- Key-value stores (like HBase, Cassandra, Redis, ...)

Basically persistent distributed hash tables

- Semi-structured/Document data stores (for XML/JSON query processing)
- Graph databases
- Data lakes (e.g., scientific data, machine learning data)

However, many lessons to be learned from database research

We see much reinvention of the wheel and similar mistakes being made as early on

Data Management Challenges (1)

- ▶ How to “represent”/“model” the data so users can easily understand, query, analyze, and update?
 - This is about the “logical” representation rather than “physical” -- we will worry about physical representation later
 - Whatever format we use needs to be easy for humans to work with, over a long period of time as data and usage evolves

As tables (like Excel)?

ID	name	dept_name	salary
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000

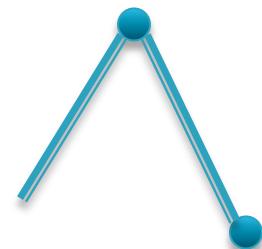
dept_name	building	budget
Comp. Sci.	Taylor	100000
Biology	Watson	90000
Elec. Eng.	Taylor	85000

JSON Documents?

```
{"instructors": [  
    {"ID": "22222",  
     "name": "Einstein",  
     "dept_name": "Physics",  
     ....  
    },  
    ....  
]
```

Graph?

Instructor: ID = 22222



Department: Physics

Data Management Challenges (1)

- ▶ How to “represent”/“model” the data so users can easily understand, query, analyze, and update?
 - This is about the “logical” representation rather than “physical” -- we will worry about physical representation later
 - Whatever format we use needs to be easy for humans to work with, over a long period of time as data and usage evolves
 - Many questions...
 - Pros and cons of different ways to do this
 - How to discuss whether our chosen data format is good or bad
 - e.g., some ways of storing data may increase conflicts in the data
 - How to deal with changes to the requirements and types of data being stored
 - Backwards compatibility is a huge problem leading to messy representations later on

Data Modeling

Data Management Challenges (2)

- ▶ How to query the data? How to process/analyze it ?
(from users'/analysts' perspective)
 - “find the stores with the maximum increase in sales in last month”
 - We can't expect the users to write Java programs
 - “how much time from here to Pittsburgh if I start at 2pm ?”
 - Data is there (GPS, live traffic data), but requires predictive capabilities
 - Increasing need to convert “information” to “knowledge”: **Data mining/Machine Learning**
 - Predicting which TikToks/tweets to show, ...
 - Needs to be easy to use, but powerful enough for what needs to be done
 - Otherwise users will simply take out all the data and write their own programs

Query Languages and
Programming Abstractions

Data Management Challenges (3)

- ▶ How to query the data? How to process/analyze it ?
(from implementation perspective)
 - How to handle the increasing scale of the data
 - How to support more and more complex analytics/machine learning over the data
 - How to process the data in real-time, as it is arriving into the system
 - How to find the best way to execute a query with least resources/time
 - How to utilize large numbers of machines to execute queries more efficiently
 - ...

Database Query Engines
Big Data Systems

Data Management Challenges (4)

- ▶ Modifying/updating the data given:
 - Many users simultaneously updating the data...
 - While the data is (potentially) replicated across the globe...
 - And any server or the network may go down at any point.

Update 1 Update 2

```
a = a + 10    inc = a * 0.15
b = b - 10    a = a + inc
save a & b    b = b - inc
               save a & b
```

If allowed to
run
concurrently

```
a = a + 10
inc = a * 0.15
a = a + inc
b = b - 10
save a & b
b = b - inc
save a & b
```

Doesn't happen if forced to run one after the other, but unacceptable performance

Data Management Challenges (4)

- ▶ Modifying/updating the data given:
 - Many users simultaneously updating the data...
 - While the data is (potentially) replicated across the globe...
 - And any server or the network may go down at any point.

Update

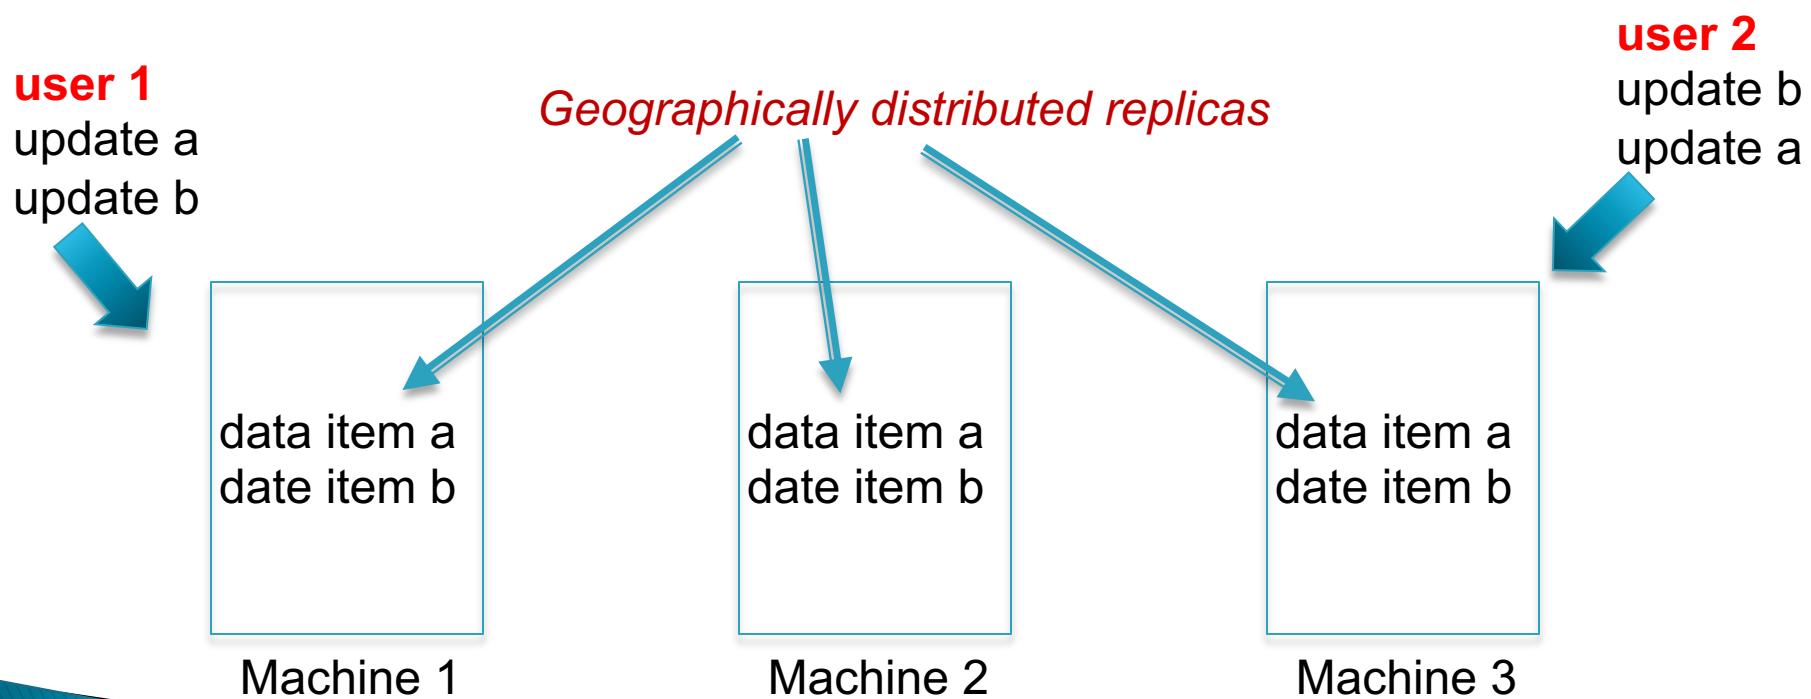
```
a = a + 10  
b = b - 10  
save a  
===== System crashes  
save b
```

Need to figure out how to handle partially finished transactions

More complications if we do update-in-place (i.e., $a = a + 10$ updates the original value)

Data Management Challenges (4)

- ▶ Modifying/updating the data given:
 - Many users simultaneously updating the data...
 - While the data is (potentially) replicated across the globe...
 - And any server or the network may go down at any point.



Data Management Challenges (4)

- ▶ Modifying/updating the data given:
 - Many users simultaneously updating the data...
 - While the data is (potentially) replicated across the globe...
 - And any server or the network may go down at any point.
- ▶ From Users' perspective:
 - What kinds of guarantees can be provided to the users about the end result, and about the order in which updates are made?
 - What if the users are okay with weaker guarantees?
- ▶ From Implementation perspective:
 - How to provide the guarantees efficiently?

Transactions

Data Management Challenges (5-∞)

- ▶ Data Integration, Transformation, and Cleaning
 - Large amounts of time spent in just getting data ready for analysis, by some accounts 90-95%
 - Real-world data has a lot of noise, missing values, etc...
- ▶ Security
 - Access control, use of encryption, ...
- ▶ Privacy
 - How to allow users to query data without revealing it
 - How to enable public data releases (e.g., by Census Bureau)
- ▶ ...

Database System

- ▶ A DBMS is a software system designed to store, manage, facilitate access to databases
 - Typically uses a specific *data model*, and
 - Supports some level of *physical and logical data independence*
- ▶ Provides:
 - Data Definition Language (DDL)
 - For defining and modifying the schemas
 - Data Manipulation Language (DML)
 - For retrieving, modifying, analyzing the data itself
 - Guarantees about correctness in presence of failures and concurrency, data semantics etc.

Example: Relational DBMS and SQL

- ▶ **SQL** (sequel): Structured Query Language

- ▶ **Data definition (DDL)**

- **create table** *instructor* (
 ID **char(5)**,
 name **varchar(20)**,
 dept_name **varchar(20)**,
 salary **numeric(8,2)**)

- ▶ **Data manipulation (DML)**

- Example: Find the name of the instructor with ID 22222

```
select name  
from instructor  
where instructor.ID = '22222'
```

Example: MongoDB and MongoQL

► Data definition (DDL)

- Tuple/Row → Document; Relation/Table → Collection

```
db.createCollection("people")
```

- Also, “insert” implicitly creates collections

```
db.people.insertOne(  
  { user_id: "abc123",  
    age: 55,  
    status: "A" } )
```

► Data manipulation (DML)

- Example: Find all users with status = “A” and return user_id and status, but don’t return _id

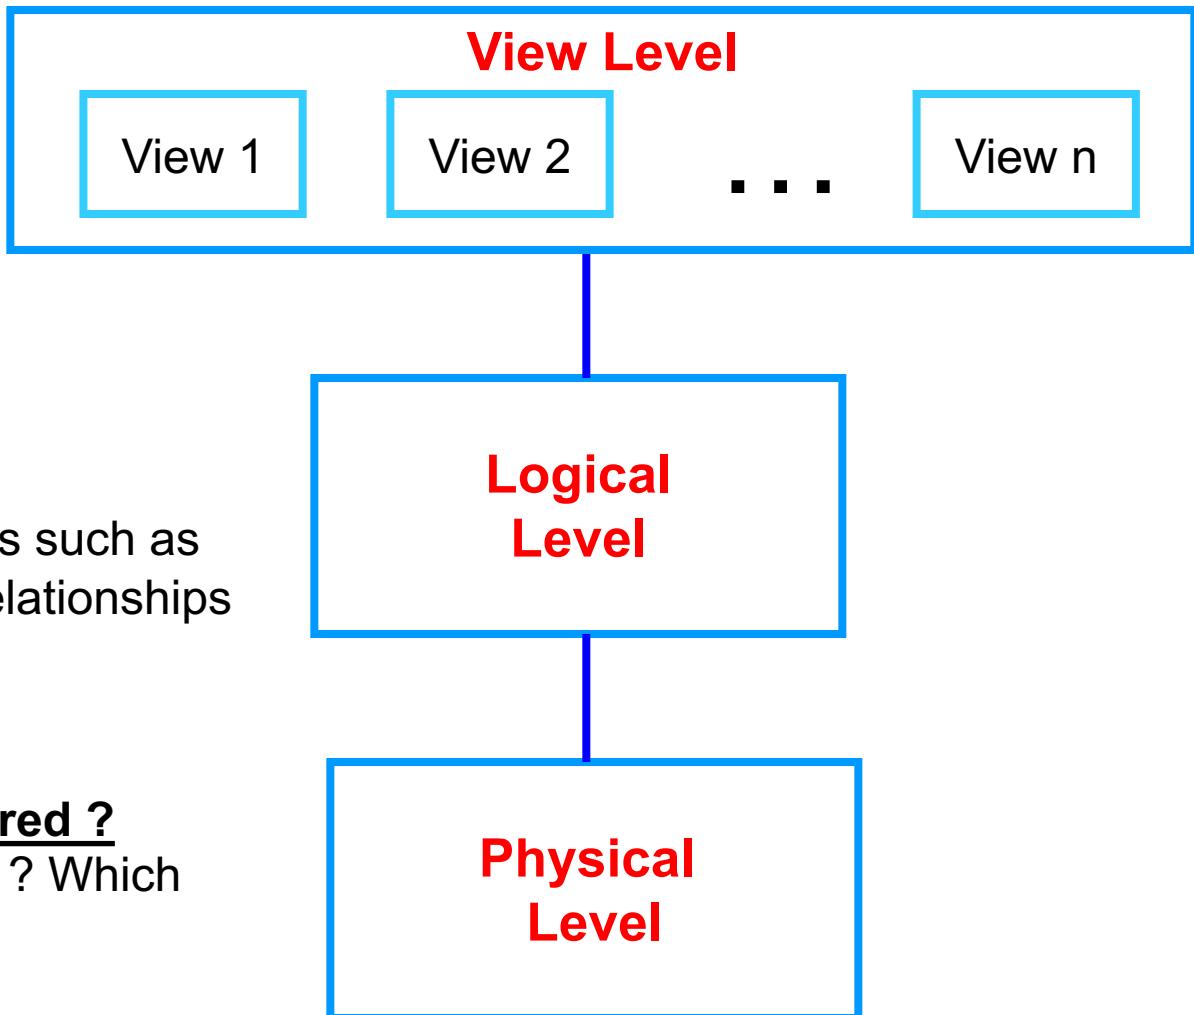
```
db.people.find(  
  { status: "A" },  
  { user_id: 1, status: 1, _id: 0 } )
```

Data Abstraction

- ▶ Probably *the* most important purpose of a DBMS
- ▶ Goal: Hiding *low-level details* from the users of the system
 - Alternatively: the principle that
 - *applications and users should be insulated from how data is structured and stored*
 - Also called *data independence*
- ▶ Through use of *logical abstractions*

Data Abstraction

What data users and application programs see ?



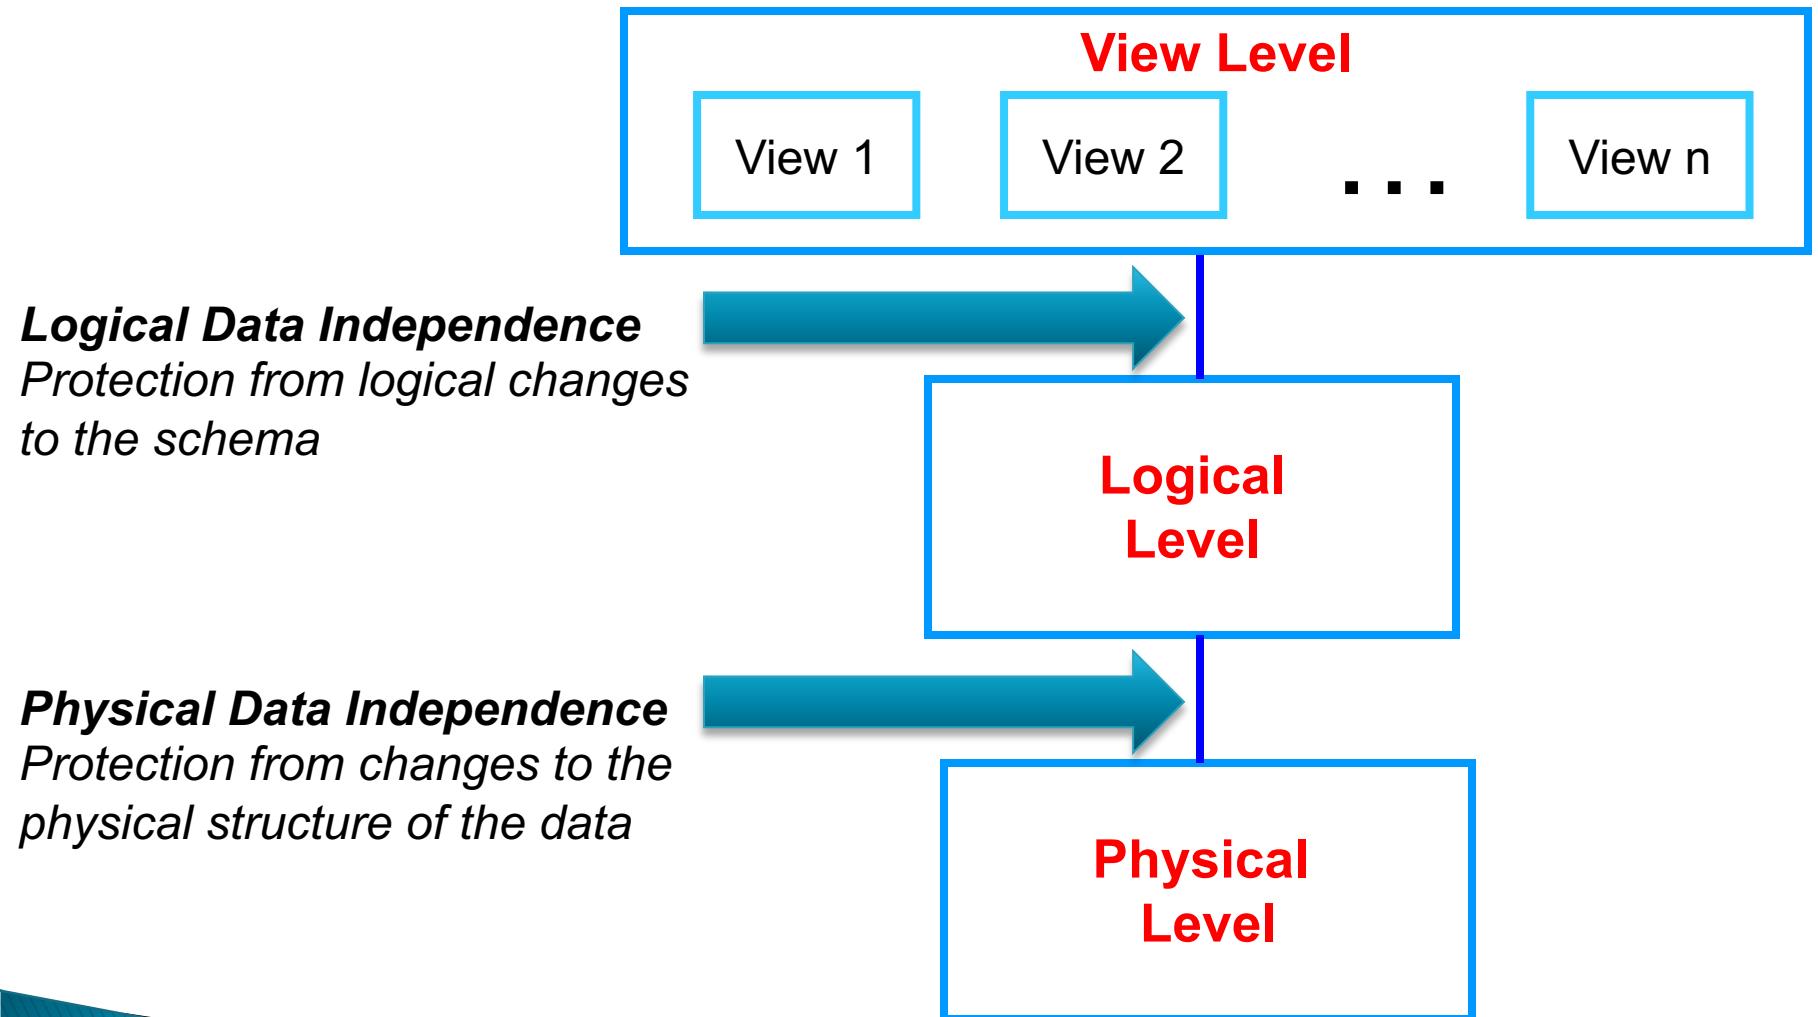
What data is stored ?

describe data properties such as data semantics, data relationships

How data is actually stored ?

e.g. are we using disks ? Which file system ?

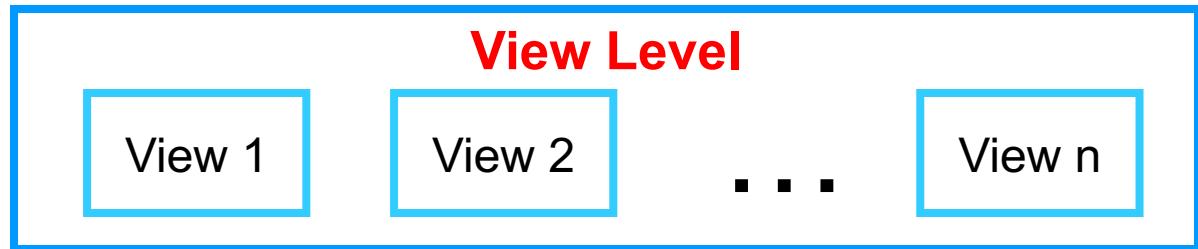
Data Abstraction



Data Abstractions: Example

A View Schema

course_info(#registered, ...)

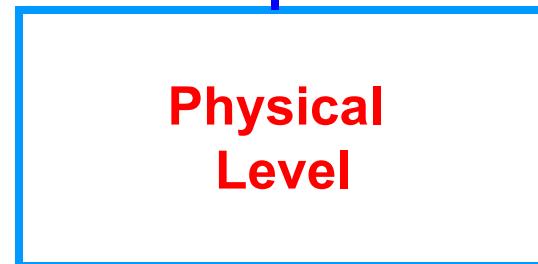


Logical Schema

students(sid, name, major, ...)

courses(cid, name, ...)

enrolled(sid, cid, ...)



Physical Schema

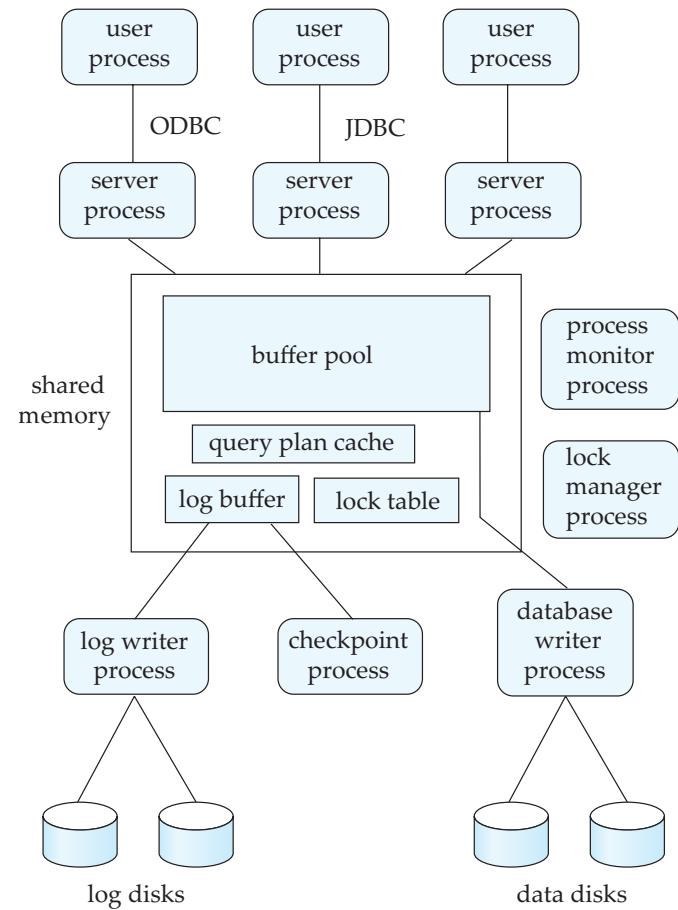
all students in one file ordered by sid

courses split into multiple files by colleges

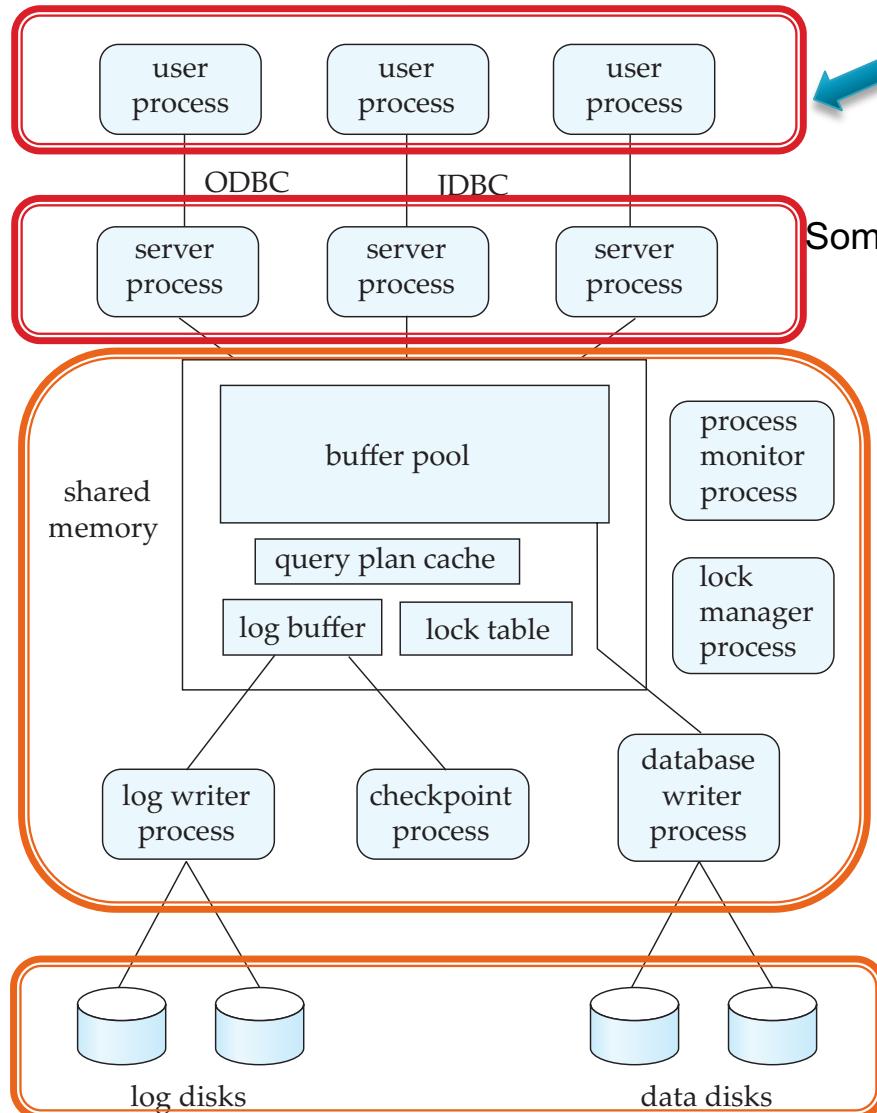
Database Architecture: Pre-2000's

- ▶ All data was typically in hard disks or arrays of hard disks
- ▶ RAM (Memory) was never enough
 - So always had to worry about what was in memory vs not
- ▶ Almost no real “distributed” execution
 - Different from “parallel”, i.e., on co-located clusters of computers
- ▶ Relatively well-understood use cases
 - Report generation
 - Interactive data analysis and exploration
 - Supporting transactions

From Chapter 20



Traditional RDBMS Architecture



Clients may be anywhere – e.g., ATMs, desktops, laptops, web apps etc.

Talk to the database using standard protocols like JDBC/ODBC, SOAP, or REST (today), or proprietary protocols

Some sort of load balancer or intake mechanism

Typical components in a database system: some for queries, some for transactions

Maybe on a single physical computer or a cluster connected by a fast network

Data Storage Systems:

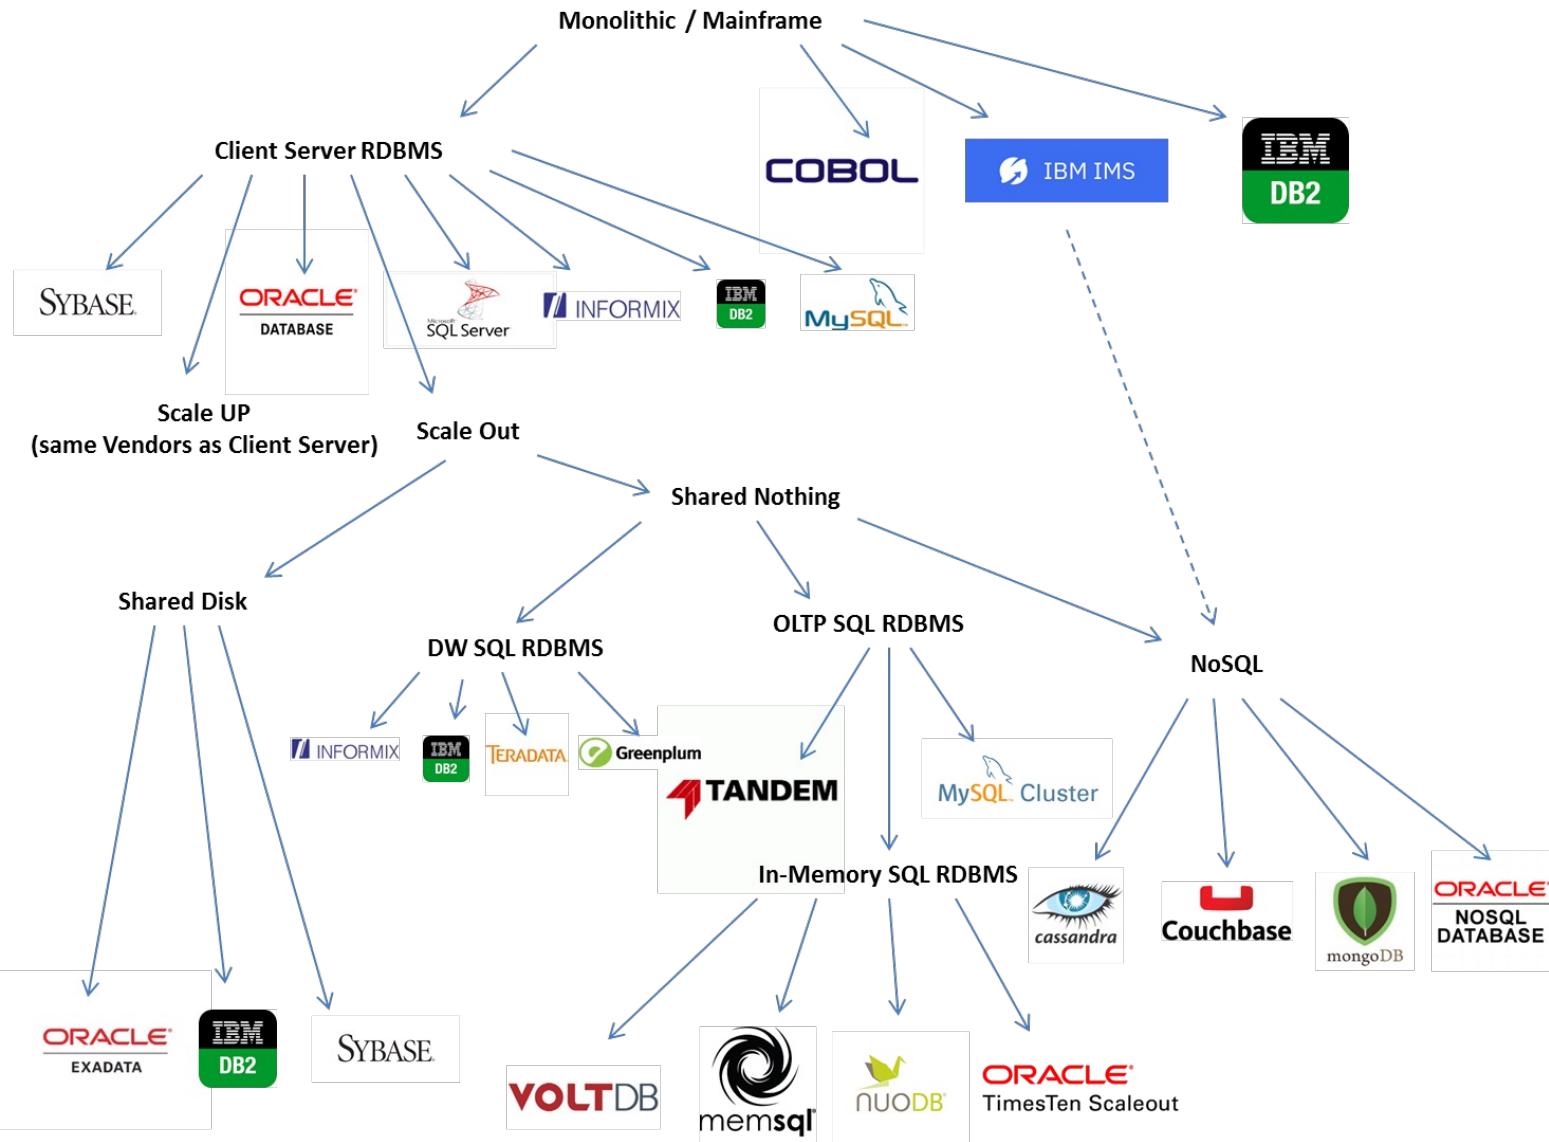
- (1) Punch cards (long time ago)
- (2) Hard disks (still prevalent)
- (3) SSDs

Need “redundancy” and “fault-tolerance”
Data once stored should always be there

RAID = Redundant Array of Independent Disks

Database Architecture : Today

- ▶ Much more diversity in the architectures that we see
 - More modern hardware architectures
 - Massively parallel computers
 - SSDs
 - Massive amounts of RAM – often don't need to worry about data fitting in memory
 - Much faster networks, even over a wide area
 - Virtualization and Containerization
 - Cloud Computing
 - As a result: Data and execution typically distributed all over the place
- ▶ Much more diversity in data processing applications
 - Much more non-relational data (images, text, video)
 - Data Analytics/Machine learning more common use-cases
- ▶ Much more diversity in “data models”
 - Document data models (JSON, XML), Key-value data model, Graph data model, RDF



From: <https://blogs.oracle.com/timesten/the-evolution-of-db-architectures>
 (Oracle-focused)

Data Warehouses

For: Large-scale data processing (TBs to PBs)

Parallel architectures (lots of co-located computers)

SQL and Reporting

No transactions

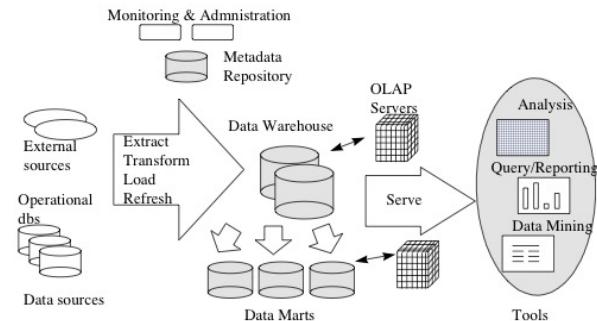


Figure 1. Data Warehousing Architecture

In-memory OLTP (on-line transaction processing)

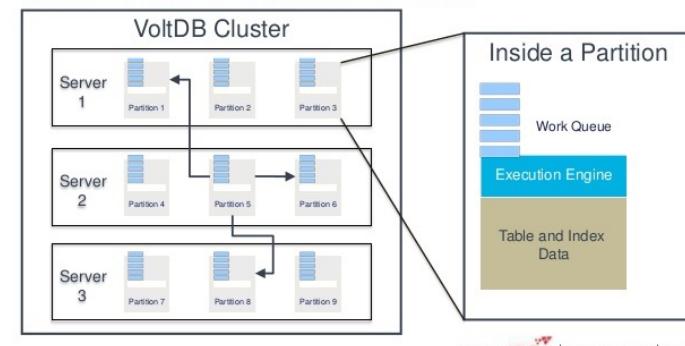
For: Extremely fast transactions

Many-core or parallel architectures

Very limited SQL – mostly focused on “writes”

Typically assume data fits in memory across servers

VOLTDB: A BEAUTIFUL ARCHITECTURE



Highly available, distributed OLTP

For: Distributed scenarios where clients are all over the world

Focus on “consistency” – how to make sure all users see the same data

Limited SQL – mostly focused on “writes”

Considerations of memory vs disk less important



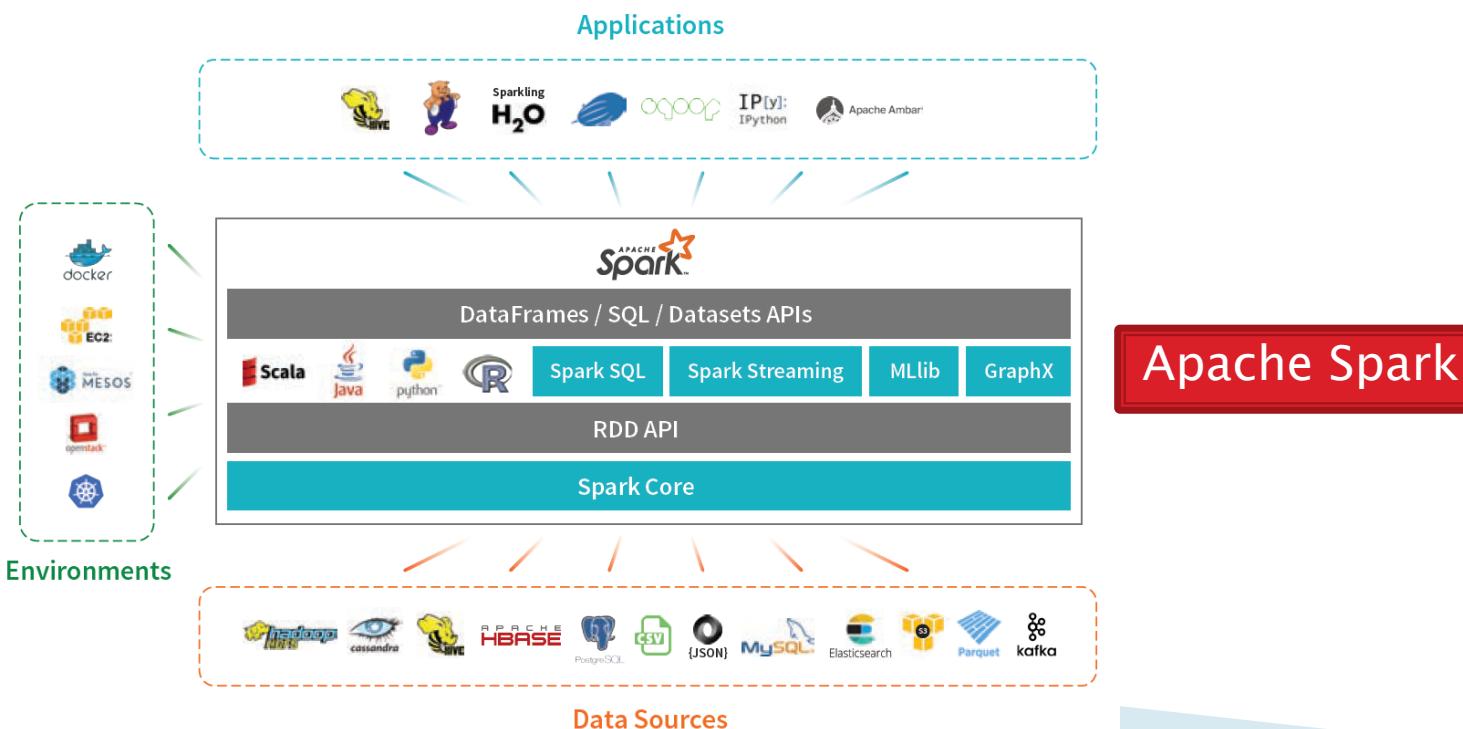
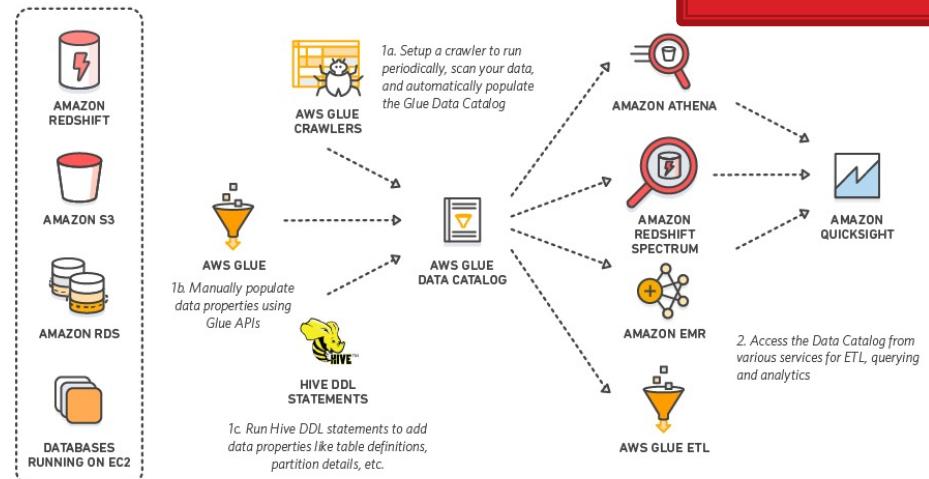
Extract-Transform-Load Systems, or Map-Reduce, or Big Data Frameworks

For: Large-scale, “ad hoc” data analysis

Mix of parallel and distributed architectures

Data usually coming from many different sources

Mix of SQL, Machine Learning, and ad hoc tasks (e.g., do image analysis, followed by SQL)



Okay...

- ▶ Key takeaway: Modern data architectures are all over the place
 - We haven't talked about NoSQL (MongoDB, etc.), Machine Learning, "Streaming"...
- ▶ Fundamentals haven't changed that much though
 - We are still either:
 - Going from some "input datasets" to an "output dataset" (queries/analytics)
 - Modifying data (transactions)
 - SQL is still very common, albeit often disguised
 - Spark RDD operations map nicely to SQL joins and aggregates (unified now)
 - MongoDB lookups, filters, and aggregates map to joins, selects, and aggregates in SQL
- ▶ But "performance trade-offs" are all over the place now
 - 30 years ago, we worried a lot about hard disks and things fitting in memory
 - Today, focus more on networks
- ▶ Focus has shifted to other aspects of data processing pipelines
 - Analytics/Machine learning, data cleaning, statistics

What we will cover...

- ▶ Data Modeling
 - Data represented in tabular forms (Relational)
 - JSON/Document Data Model (MongoDB)
 - Entity-relationship Model for Schema Design
 - Normalization theory to formally define a "good" schema
 - We will discuss a more primitive data model as well (Apache Spark)
- ▶ Query Languages and Programming Abstractions
 - SQL, MongoDB Query Language, Apache Spark MapReduce-like Framework, Some others
- ▶ Query Engines and Big Data Frameworks
 - Considerations in efficient implementations
- ▶ Transactions
 - How to reason about consistency and durability guarantees (ACID)
 - Basics of how to achieve those guarantees in centralized and distributed settings

What we will cover...

- ▶ This course is NOT about learning how to write SQL queries
 - Although we will practice that a fair bit...
- ▶ It is about learning the foundations of how data management systems are built and why they are built that way, so you:
 - Can think about the design decisions more systematically (e.g., through use of normalization theory and E-R modeling)
 - Understand the role of data modeling and how different models impact design decisions
 - Develop a deeper understanding of the differences and similarities between different types of database systems (e.g., PostgreSQL vs MongoDB)
 - Are equipped (to some extent) to design next-generation data systems
 - Are able to reason about performance of the queries/tasks

Structure of the Course

- ▶ Introduction
 - Motivation, data abstraction, common data systems architectures today
- ▶ Relational Model + SQL (**Three Programming assignments**)
 - How to model and query data using SQL
 - How to update data using transactions and considerations therein
- ▶ Schema Design: Entity-relationship Models and Normalization (**One assignment**)
 - How to create a database schema, and how to ensure it is “good”
- ▶ NoSQL (somewhat of a misnomer) (**Two Programming assignments**)
 - Document, key-value, and graph data models
 - MongoDB and its Query Language
 - Map-reduce Model and Apache Spark
- ▶ Implementation Issues
 - Different types of storage, and how to ensure reliability in presence of failures
 - Indexes for faster retrieval of data
 - How an SQL query is processed and optimized
 - How to do concurrent updates correctly
 - How to ensure consistency in presence of failures